

# Report

Kubernetes, Docker Swarm & Microservices.

**INFO 7390**

**Advance Data Science & Architecture**

Professor: Srikanth Krishnamurthy

By: Team 5

Akash Jagtap

Jerin Rajan


Nitin Prince Reuben

By this report, we can learn about online hosting applications like Kubernetes, docker swarm and basic concepts of microservices.


## Kubernetes.

Kubernetes, the word is derived from Greek which means ‘helmsman’ or pilot of a ship. It is an open source project that was started by Google and derived from Borg, which is used inside Google for several years now, for container management, automating deployment and scaling. Google's Borg system is a cluster manager that runs hundreds of thousands of jobs, from many thousands of different applications, across a number of clusters each with up to tens of thousands of machines. It achieves high utilization by combining admission control, efficient task-packing, over-commitment, and machine sharing with process-level performance isolation. It supports high-availability applications with runtime features that minimize fault-recovery time, and scheduling policies that reduce the probability of correlated failures. Borg simplifies life for its users by offering a declarative job specification language, name service integration, real-time job monitoring, and tools to analyze and simulate system behavior. Currently, it is hosted by Cloud Native Computing Foundation(CNCF). There are lot of things along with kubernetes which is being hosted by CNCF for various purposes. Major ones are Prometheus for monitoring and CNI for networking API. They are trying to gather as many software/applications as possible which can provide smooth transition for any company towards cloud business. You can learn more about these software from <https://www.cncf.io/>.

The old way: Applications on host



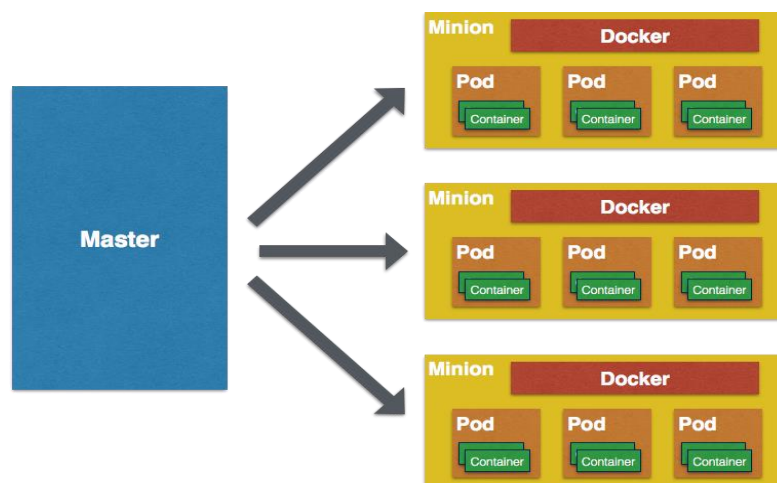
The new way: Deploy containers




By this picture, we can understand the difference between older and new technique of deploying the applications. Earlier, the containers to which the applications are being deployed has some libraries already present in it. That had few outcomes. Not all applications have same requirements. Some preferred older libraries and some preferred new. This created lots of problems. This made people change the deployment technique. Now the libraries which are required by the applications are being deployed with it. This decreases the problem of compatibility and increased the performance and speed. This also made upgradations easier.

Kubernetes have lots of features. Some of them are as follows:

- Does not limit the types of application supported – It supports wide varieties of application and this make it preferable over other deployment application
- Load Balancing – if the traffic to the network is more, it helps has to balance the load easily.
- Rolling update or rollback – it maintains logs. So, it is easy to roll back to a previous version or to make an update easy.
- Auto – scalability
- Using pods for grouping container
- API management – It provides lots API to work upon. All these API will be disabled at the beginning. We can start using the ones which we need by enabling them. And can later disable them after we are done using it as google charges money for using these API depending upon the time of usage. So, we don't want to keep unused API enabled. Also, google thinks that by disabling the APIs, it can increase security as people can attack your account to steal the code or cause a problem to your application by using the APIs which we are not using.




The above picture describes the architecture of Kubernetes. For a clearer, understanding, we would check next image.



Kubernetes is made up of clusters. Each cluster has a master surrounded by nodes.


## Node overview



Each node consists of lots of pods inside which the application is deployed and every pod have it's set of libraries already present in it.

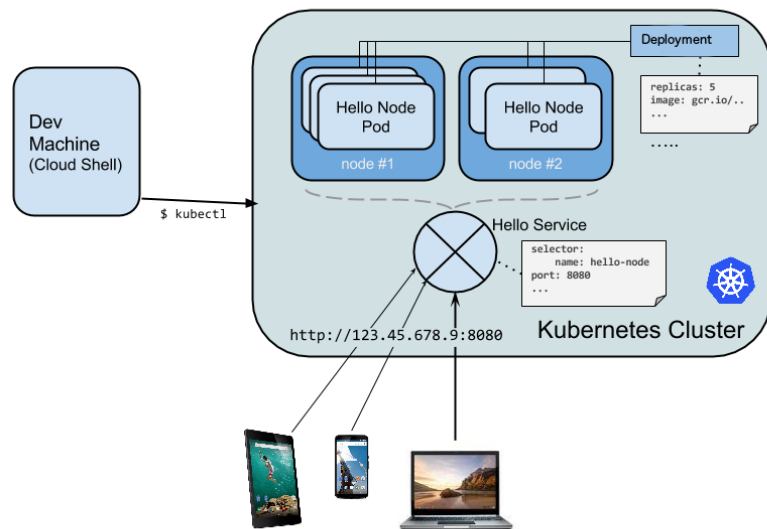

## Pods overview

---



The basic scheduling unit in Kubernetes is called a "pod". It adds a higher level of abstraction to containerized components. A pod consists of one or more containers that are guaranteed to be co-located on the host machine and can share resources. Each pod in Kubernetes is assigned a unique (within the cluster) IP address, which allows applications to use ports without the risk of conflict.] A pod can define a volume, such as a local disk directory or a network disk, and expose it to the containers in the pod. Pods can be managed manually through the Kubernetes API, or their management can be delegated to a controller.

A Kubernetes service is a set of pods that work together, such as one tier of a multi-tier application. The set of pods that constitute a service are defined by a label selector.] Kubernetes provides service discovery and request routing by assigning a stable IP address and DNS name to the service, and load balances traffic in a round-robin manner to network connections of that IP address among the pods matching the selector (even as failures cause the pods to move from machine to machine). By default a service is exposed inside a cluster (e.g. back end pods might be grouped into a service, with requests from the front-end pods load-balanced among them), but a service can also be exposed outside a cluster (e.g. for clients to reach frontend pods).



With the above picture, we can understand the business prospective of Kubernetes. Through our developer machine, we transfer the code. *Kubectl* is the function which helps google SDK to communicate with kubernetes. With the help of auto-scaling, the code is distributed and we are provided with an IP address with which we can access the application in our devices.

## Deploying a Simple Hello World App on Kubernetes.

Technology used are: Docker, Python, Flask, Kubernetes, Google Cloud Platform

1. Create a Google Cloud Account
2. Install gcloud SDK (client cli)
3. After installing SDK there are two ways to setups kubernetes
  - a. Minicube – Single Node Kubernetes
    - i. For installation – [https://www.youtube.com/watch?v=\\_vHTaIJm9uY](https://www.youtube.com/watch?v=_vHTaIJm9uY)
4. Kubernetes multi-node cluster can be performed on AWS or GCP.  
We would be proceeding with GCP i.e. Google Cloud Platform.
5. Commands to execute are :
  - i. gcloud components list

```
frivolouspantheon jerry # gcloud components list
Your current Cloud SDK version is: 188.0.1
The latest available version is: 189.0.0
```

Components			
Status	Name	ID	Size
Update Available	BigQuery Command Line Tool	bq	< 1 MiB
Update Available	Cloud SDK Core Libraries	core	7.2 MiB
Update Available	Cloud Storage Command Line Tool	gsutil	3.3 MiB
Not Installed	App Engine Go Extensions	app-engine-go	98.1 MiB
Not Installed	Cloud Bigtable Command Line Tool	cbt	4.5 MiB
Not Installed	Cloud Bigtable Emulator	bigtable	3.7 MiB
Not Installed	Cloud Datalab Command Line Tool	datalab	< 1 MiB
Not Installed	Cloud Datastore Emulator	cloud-datastore-emulator	17.9 MiB
Not Installed	Cloud Datastore Emulator (Legacy)	gcd-emulator	38.1 MiB
Not Installed	Cloud Pub/Sub Emulator	pubsub-emulator	33.4 MiB
Not Installed	Emulator Reverse Proxy	emulator-reverse-proxy	14.5 MiB
Not Installed	gcloud Alpha Commands	alpha	< 1 MiB
Not Installed	gcloud Beta Commands	beta	< 1 MiB
Not Installed	gcloud app Java Extensions	app-engine-java	118.8 MiB
Not Installed	gcloud app PHP Extensions	app-engine-php	
Not Installed	gcloud app Python Extensions	app-engine-python	6.1 MiB
Not Installed	gcloud app Python Extensions (Extra Libraries)	app-engine-python-extras	27.8 MiB
Installed	Google Container Local Builder	container-builder-local	3.8 MiB
Installed	Google Container Registry's Docker credential helper	docker-credential-gcr	2.8 MiB
Installed	kubectl	kubectl	12.3 MiB

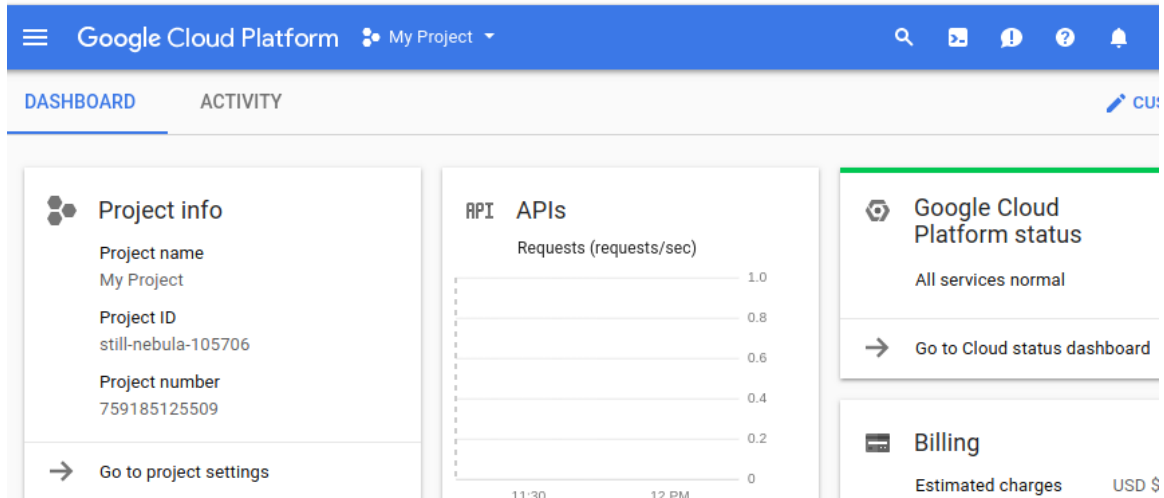
```
To install or remove components at your current SDK version [188.0.1], run:
$ gcloud components install COMPONENT_ID
$ gcloud components remove COMPONENT_ID

To update your SDK installation to the latest version [189.0.0], run:
$ gcloud components update
```

- ii. If kubectl is not installed then install it using the command  
gcloud component install kubectl

Ref:<https://cloud.google.com/sdk/gcloud/reference/components/>

6. Login with CLI:
  - a. gcloud auth login
  - b. Enter your username and password
  - c. Go to your gcloud console and create a project.



Next, Click + sign And assign your project name

**Select**

Search projects and folders

Recent All

Name	ID
kubernetes-app	kubernetes-app-195105

7. Goto the CLI
  - a. `gcloud config set project[PROJECT_ID]`
  - b. `gcloud config set compute/zone [COMPUTE_ZONE]`
  - OR
  - c. `gcloud init`
    - i. Enter Option 1.
    - ii. Login with your username
    - iii. Select the project you just created.
    - iv. Select the configure zone.
8. Create Kubernetes Cluster

A cluster consists of at least one cluster master machine and multiple machines called nodes. Nodes are Compute Engine Virtual Machine (VM) instances that run the Kubernetes processes necessary to make them part of the cluster. You deploy applications to clusters and the applications run on the nodes

Command to create cluster is : `gcloud container clusters create kubecuster101`



Google Cloud Platform **kubernetes-app**

**Kubernetes clusters** [CREATE CLUSTER](#) [REFRESH](#) [DELETE](#)

Filter by label or name

**Kubernetes clusters**

Name	Location	Cluster size	Total cores	Total memory	Notifications	Labels
<input checked="" type="checkbox"/> kubercluster101	us-east1-b	3	3 vCPUs	11.25 GB		

[Connect](#) [Edit](#) [Delete](#)

**Instance groups** [CREATE INSTANCE GROUP](#) [REFRESH](#) [EDIT](#) [DELETE](#)

Filter instance groups Columns

Name	Zone	Creation time	Instances	Template	Recommendation	Autoscaling	In use by
<input checked="" type="checkbox"/> gke-kubercluster101-default-pool-68aaf079-grp	us-east1-b	Feb 19, 2018, 12:09:22 PM	3	gke-kubercluster101-default-pool-68aaf079		Off	

Ref: <https://kubernetes.io/docs/setup/pick-right-solution/#hosted-solutions>

9. Push Docker images to Container Registry
  - a. Create a Docker File

# Import Python runtime and set up working directory

```
FROM python:2.7-slim
WORKDIR /app
ADD . /app
# Install any necessary dependencies
RUN pip install -r requirements.txt
# Open port 80 for serving the webpage
EXPOSE 80
# Run app.py when the container launches
CMD ["python", "app.py"]
```

- b. Create requirements.txt << Flask
- c. Create app.py

```
from flask import Flask

import os
import socket
app = Flask(__name__)
@app.route("/")
def hello():
    html = "<h3>Hello, World!</h3>"
    return html
if __name__ == "__main__":
    app.run(host='0.0.0.0', port=80)
```

```
Kubernetes # vi Dockerfile
Kubernetes # vi requirements.txt
Kubernetes # vi app.py
Kubernetes #
```

- d. Build docker images

```
Frivolouspantheon Kubernetes # docker build -t quickstart-image .
Sending build context to Docker daemon 4.096 kB
Step 1/6 : FROM python:2.7-slim
--> 52ad41c7aead
Step 2/6 : WORKDIR /app
--> 89c9072e73dc
Removing intermediate container ed19bad66666
Step 3/6 : ADD . /app
--> 639567118ceb
Removing intermediate container edf864f1ed60
Step 4/6 : RUN pip install -r requirements.txt
--> Running in a20e7b1459d6
Collecting Flask (from -r requirements.txt (line 1))
  Downloading Flask-0.12.2-py2.py3-none-any.whl (83kB)
Collecting itsdangerous==0.21 (from Flask->-r requirements.txt (line 1))
  Downloading itsdangerous-0.24.tar.gz (46kB)
Collecting Jinja2==2.4 (from Flask->-r requirements.txt (line 1))
  Downloading Jinja2-2.10-py2.py3-none-any.whl (126kB)
Collecting Werkzeug==0.7 (from Flask->-r requirements.txt (line 1))
  Downloading Werkzeug-0.14.1-py2.py3-none-any.whl (322kB)
Collecting click==2.0 (from Flask->-r requirements.txt (line 1))
  Downloading click-6.7-py2.py3-none-any.whl (71kB)
Collecting MarkupSafe==0.23 (from Jinja2==2.4->Flask->-r requirements.txt (line 1))
  Downloading MarkupSafe-1.0.tar.gz
Building wheels for collected packages: itsdangerous, MarkupSafe
  Running setup.py bdist_wheel for itsdangerous: started
  Running setup.py bdist_wheel for itsdangerous: finished with status 'done'
  Stored in directory: /root/.cache/pip/wheels/fc/a8/66/24d655233c757e178d45dea2de22a04c6d92766abfb741129a
  Running setup.py bdist_wheel for MarkupSafe: started
  Running setup.py bdist_wheel for MarkupSafe: finished with status 'done'
  Stored in directory: /root/.cache/pip/wheels/88/a7/30/e39a54a87bcbe25308fa3ca64e8ddc75d9b3e5afa21ee32d57
Successfully built itsdangerous MarkupSafe
Installing collected packages: itsdangerous, MarkupSafe, Jinja2, Werkzeug, click, Flask
Successfully installed Flask-0.12.2 Jinja2-2.10 MarkupSafe-1.0 Werkzeug-0.14.1 click-6.7 itsdangerous-0.24
--> 11b3e9be36e2
Removing intermediate container a20e7b1459d6
Step 5/6 : EXPOSE 80
--> Running in 06195fd43105
--> b9cb682c42f1
Removing intermediate container 06195fd43105
Step 6/6 : CMD python app.py
--> Running in 90bbcc9d4ec5
--> 745503856a7f
Removing intermediate container 90bbcc9d4ec5
Successfully built 745503856a7f
```

e. Tagging your image

Before you push your Docker image, you need to tag it with its registry name. Tagging your Docker image with a registry configuration the docker push command to push the image to a specific location.

The registry name format is below:

```
[HOSTNAME]/[PROJECT-ID]/[IMAGE]
```

where

- **[HOSTNAME]** is the **gcr.io** hostname
- **[PROJECT-ID]** is your Google Cloud Platform Console project ID
- **[IMAGE]** is your image's name

To tag your Docker image for Container Registry, run the following command:

```
docker tag [IMAGE] [HOSTNAME]/[PROJECT-ID]/[IMAGE]
```

For example

```
docker tag quickstart-image gcr.io/my-project/quickstart-image
```

```
Kubernetes # docker tag quickstart-image gcr.io/kubernetes-app-195105/quickstart-image
Kubernetes #
```



f. Pushing your image

# Pushing your image

To push your Docker image to Container Registry, run the following command:

```
gcloud docker -- push [HOSTNAME]/[PROJECT-ID]/[IMAGE]
```


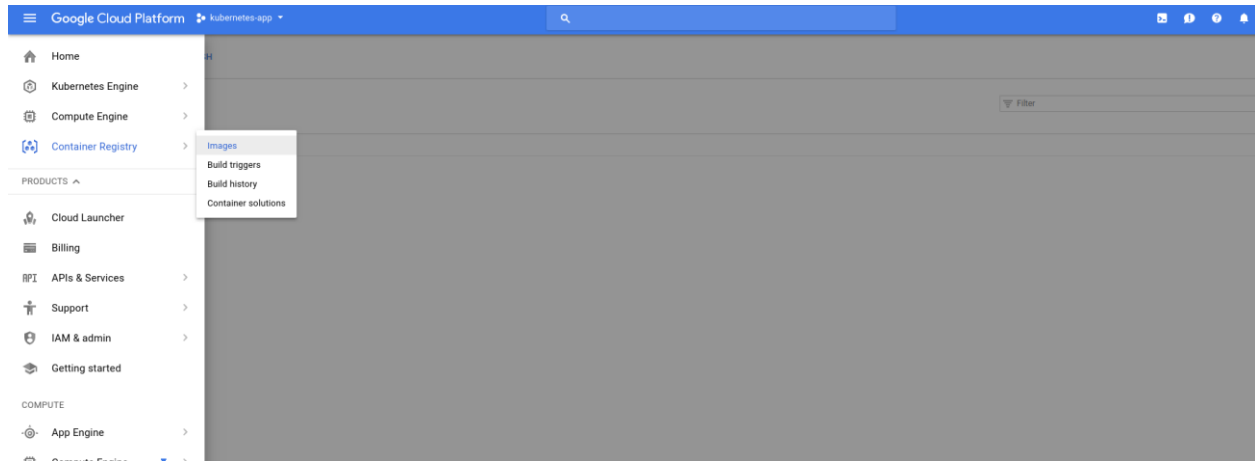
For example:

```
gcloud docker -- push gcr.io/my-project/quickstart-image
```

```
frivolouspantheon Kubernetes # docker build -t quickstart-image .
Sending build context to Docker daemon 4.096 kB
Step 1/6 : FROM python:2.7-slim
--> 52ad41c7aea4
Step 2/6 : WORKDIR /app
--> 89c9072e73dc
Removing intermediate container ed19bad66666
Step 3/6 : ADD . /app
--> 639567118ceb
Removing intermediate container edf864fled60
Step 4/6 : RUN pip install -r requirements.txt
--> Running in a20e7b1459d6
Collecting Flask (from -r requirements.txt (line 1))
  Downloading Flask-0.12.2-py2.py3-none-any.whl (83kB)
Collecting itsdangerous>=0.21 (from Flask->-r requirements.txt (line 1))
  Downloading itsdangerous-0.24.tar.gz (46kB)
Collecting Jinja2>=2.4 (from Flask->-r requirements.txt (line 1))
  Downloading Jinja2-2.10-py2.py3-none-any.whl (126kB)
Collecting Werkzeug>=0.7 (from Flask->-r requirements.txt (line 1))
  Downloading Werkzeug-0.14.1-py2.py3-none-any.whl (322kB)
Collecting click>=2.0 (from Flask->-r requirements.txt (line 1))
  Downloading click-6.7-py2.py3-none-any.whl (71kB)
Collecting MarkupSafe>=0.23 (from Jinja2>=2.4->Flask->-r requirements.txt (line 1))
  Downloading MarkupSafe-1.0.tar.gz
Building wheels for collected packages: itsdangerous, MarkupSafe
  Running setup.py bdist_wheel for itsdangerous: started
  Running setup.py bdist_wheel for itsdangerous: finished with status 'done'
  Stored in directory: /root/.cache/pip/wheels/fc/a8/66/24d655233c757e178d45dea2de22a04c6d92766abfb741129a
  Running setup.py bdist_wheel for MarkupSafe: started
  Running setup.py bdist_wheel for MarkupSafe: finished with status 'done'
  Stored in directory: /root/.cache/pip/wheels/88/a7/30/e39a54a87bcb25308fa3ca64e8ddc75d9b3e5afa21ee32d57
Successfully built itsdangerous MarkupSafe
Installing collected packages: itsdangerous, MarkupSafe, Jinja2, Werkzeug, click, Flask
Successfully installed Flask-0.12.2 Jinja2-2.10 MarkupSafe-1.0 Werkzeug-0.14.1 click-6.7 itsdangerous-0.24
--> 11b3e9be36e2
Removing intermediate container a20e7b1459d6
Step 5/6 : EXPOSE 80
--> Running in 06195fd43105
--> b9cb682c42f1
Removing intermediate container 06195fd43105
Step 6/6 : CMD python app.py
--> Running in 90bbcc9d4ec5
--> 745503856a7f
Removing intermediate container 90bbcc9d4ec5
Successfully built 745503856a7f
frivolouspantheon Kubernetes # docker tag quickstart-image gcr.io/kubernetes-app-195105/quickstart-image
frivolouspantheon Kubernetes # docker images
REPOSITORY                                TAG                IMAGE ID           CREATED           SIZE
gcr.io/kubernetes-app-195105/quickstart-image  latest            745503856a7f      5 minutes ago    149 MB
quickstart-image                          latest            745503856a7f      5 minutes ago    149 MB
python                                    2.7-slim          52ad41c7aea4      4 days ago       139 MB
```

```
frivolouspantheon Kubernetes # gcloud docker -- push gcr.io/kubernetes-app-195105/quickstart-image
The push refers to a repository [gcr.io/kubernetes-app-195105/quickstart-image]
e12e0d7d6245: Pushed
3e0f650d74d1: Pushed
92e03a00a883: Pushed
03cd3fb86dd2: Pushed
630d02da980e: Pushed
b2f046b20847: Pushed
cf051be4e149: Layer already exists
latest: digest: sha256:e92cf93f95b1842175db06d42095856b267d25e611baf03e05c660d551b9786f size: 1787
frivolouspantheon Kubernetes #
```

Ref: <https://cloud.google.com/container-registry/docs/quickstart>



## 10. Deploying an application to the cluster

Now that we have created a cluster, we can deploy a containerized application to it.

The `kubectrl run` creates a new deployment name `qkimage1`. The deployment's pod runs the `quickstart` images in its container.



```
frivolouspantheon@Kubernetes: ~ $ kubectl run qkimage1 --image gcr.io/kubernetes-app-195105/quickstart-image --port 80
deployment "qkimage1" created
frivolouspantheon@Kubernetes: ~ $ kubectl expose deployment qkimage1 --type "LoadBalancer"
service "qkimage1" exposed
frivolouspantheon@Kubernetes: ~ $ kubectl get services qkimage1
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
qkimage1  LoadBalancer  10.59.251.37      <pending>         80:30962/TCP      6s
frivolouspantheon@Kubernetes: ~ $ kubectl get services qkimage1
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
qkimage1  LoadBalancer  10.59.251.37      <pending>         80:30962/TCP      12s
frivolouspantheon@Kubernetes: ~ $ kubectl get services qkimage1
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
qkimage1  LoadBalancer  10.59.251.37      <pending>         80:30962/TCP      14s
frivolouspantheon@Kubernetes: ~ $ kubectl get services qkimage1
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
qkimage1  LoadBalancer  10.59.251.37      <pending>         80:30962/TCP      15s
frivolouspantheon@Kubernetes: ~ $ kubectl get services qkimage1
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
qkimage1  LoadBalancer  10.59.251.37      <pending>         80:30962/TCP      17s
frivolouspantheon@Kubernetes: ~ $ kubectl get services qkimage1
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
qkimage1  LoadBalancer  10.59.251.37      <pending>         80:30962/TCP      28s
frivolouspantheon@Kubernetes: ~ $ kubectl get services qkimage1
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
qkimage1  LoadBalancer  10.59.251.37      35.229.119.9     80:30962/TCP      1m
frivolouspantheon@Kubernetes: ~ $
```

In this command:

- `--images` specifies a container image to deploy. In this case, the command pulls the example image from Google Container Registry bucket, `gcr.io/kubernetes-app-195105/quickstart-image` indicated the specific image version to pull. If a version is not specified the latest version is used.
- `--port` specifies the port that the container exposes.

## 11. Exposing the Deployment

- After deploying the application, we need to expose it to the internet so that users can access it. We can expose the application by creating a Service, a Kubernetes resource that exposes your application to external traffic.
- Passing in the `--type "LoadBalancer"` flag creates a Compute Engine load balancer for your container.



Google Cloud Platform **kubernetes-app**

Discovery & load balancing [REFRESH](#)

Services are sets of pods with a network endpoint that can be used for discovery and load balancing. Ingresses are collections of rules for routing external HTTP(S) traffic to services.

Is system object: **False** Filter Resources

Name	Status	Service Type	Endpoints	Pods	Namespace	Cluster
qkimage	Ok	Load balancer	35.231.50.185:8080 <a href="#">L2</a>	1 / 1	default	kubercluster101
qkimage1	Ok	Load balancer	35.229.119.9:80 <a href="#">L2</a>	1 / 1	default	kubercluster101

For Load Balancing, we right now use only 1 pod to run the service as per the picture above.

The screenshot shows the 'Service Details' page for a service named 'qkimage1'. The page includes a sidebar with navigation icons and a main content area with tabs for 'Details', 'Events', and 'YAML'. A message at the top states: 'Select the Stackdriver account to see charts.' Below this, a metadata table lists details for the service, including its cluster, namespace, creation time, labels, annotations, label selector, pod count, type, and external endpoints. A 'LoadBalancer' section provides information about the cluster IP, load balancer IP, and the load balancer ID. A 'Ports' table shows the port configuration. Finally, a 'Serving pods' table lists the pods currently serving the service.

Cluster	Namespace	Created	Labels	Annotations	Label selector	Pods	Type	External endpoints
kubercluster101	default	Feb 19, 2018, 12:58:17 PM	run: qkimage1	Not set	run = qkimage1	1 current / 1 desired	LoadBalancer	35.229.119.9:80

Cluster IP	Load balancer IP	Load balancer
10.59.251.37	35.229.119.9	a76d63103159e11e8af5042010a8e00a

Port	Node Port	Target Port	Protocol
80	30962	80	TCP

Name	Status	Restarts	Created on
qkimage1-5694c94cfd-4f85	Running	0	Feb 19, 2018, 12:58:08 PM

To increase the serving pods

Goto kubernetes Engine > Workloads > Click qkimage > Edit

The screenshot shows the 'Workloads' page in the Kubernetes Engine console. It includes a sidebar with navigation icons and a main content area with a 'REFRESH' button. A description at the top states: 'Workloads are deployable units of computing that can be created and managed in a cluster.' Below this, there is a filter bar with a search input, a filter for 'Is system object : False', and a 'Filter workloads' button. A table lists the workloads, showing their names, status, type, pod count, namespace, and cluster.

Name	Status	Type	Pods	Namespace	Cluster
qkimage	OK	Deployment	1/1	default	kubercluster101
qkimage1	OK	Deployment	1/1	default	kubercluster101

[Deployment details](#)
[REFRESH](#)
[EDIT](#)
[DELETE](#)
[ACTIONS](#)
[KUBECTL](#)

**qkimage1**

[Details](#)
[Revision history](#)
[Events](#)
[YAML](#)

Select the [Stackdriver](#) account to see charts.

Cluster	kubercluster101
Namespace	default
Created	Feb 19, 2018, 12:58:08 PM
Labels	run: qkimage1
Annotations	deployment.kubernetes.io/revision: 1
Replicas	1 updated, 1 ready, 1 available, 0 unavailable
Label selector	run = qkimage1
Update strategy	Rolling update, Max unavailable: 1, Max surge: 1
Min time ready before available	0 s
Progress deadline	0 s

Pod specification

Revision	1
Labels	run: qkimage1
Restart policy	Always
Containers	qkimage1

Active revisions

Revision	Name	Status	Summary	Created on	Pods running/Pods total
1	qkimage1-5694c94cfd	OK	qkimage1: gcr.io/kubernetes-app-195105/quickstart-image	Feb 19, 2018, 12:58:08 PM	1/1

Managed pods

Revision	Name	Status	Restarts	Created on
1	qkimage1-5694c94cfd-4f85	Running	0	Feb 19, 2018, 12:58:08 PM

[Deployment details](#)
[REFRESH](#)
[EDIT](#)
[DELETE](#)
[ACTIONS](#)
[KUBECTL](#)

**qkimage1**

[Details](#)
[Revision history](#)
[Events](#)
[YAML](#)

[Save](#)
[Cancel](#)

```

1 apiVersion: extensions/v1beta1
2 kind: Deployment
3 metadata:
4   annotations:
5     deployment.kubernetes.io/revision: "1"
6   creationTimestamp: 2018-02-19T17:58:08Z
7   generation: 1
8   labels:
9     run: qkimage1
10  name: qkimage1
11  namespace: default
12  resourceVersion: "4374"
13  selfLink: /apis/extensions/v1beta1/namespaces/default/deployments/qkimage1
14  uid: 711fbb6a-159e-11e8-af50-42010a8e00ae
15 spec:
16   replicas: 3
17   selector:
18     matchLabels:
19       run: qkimage1
20   strategy:
21     rollingUpdate:
22       maxSurge: 1
23       maxUnavailable: 1
24     type: RollingUpdate
25   template:
26     metadata:
27       creationTimestamp: null
28     labels:
29       run: qkimage1
30   spec:
31     containers:
32     - image: gcr.io/kubernetes-app-195105/quickstart-image
33       imagePullPolicy: Always
34       name: qkimage1
35       ports:
36       - containerPort: 80
37         protocol: TCP
38       resources: {}
39       terminationMessagePath: /dev/termination-log
40       terminationMessagePolicy: File
41     dnsPolicy: ClusterFirst

```

Here I have changed the replicas from 1 to 3.



The immediate effect which has taken place on the kubernetes cluster node.

Google Cloud Platformkubernetes app

Deployment detailsREFRESHEDITDELETEACTIONS

qkimage1

DetailsRevision historyEventsYAML

Select the Stackdriver account to see charts.

Clusterkubercluster101

Namespacedefault

CreatedFeb 19, 2018, 12:58:08 PM

Labelsrun: qkimage1

Annotationsdeployment.kubernetes.io/revision: 1

Replicas3 updated, 3 ready, 3 available, 0 unavailable

Label selectorrun = qkimage1

Update strategyRolling update, Max unavailable: 1, Max surge: 1

Min time ready before available0 s

Progress deadline0 s

Pod specification

Revision1

Labelsrun: qkimage1

Restart policyAlways

Containersqkimage1

Active revisions

Revision	Name	Status	Summary	Created on	Pods running/Pods total
1	qkimage1-5694c94cfd	OK	qkimage1: gcr.io/kubernetes-app-195105/quickstart-image	Feb 19, 2018, 12:58:08 PM	3/3

Managed pods

Revision	Name	Status	Restarts	Created on
1	qkimage1-5694c94cfd-4f85	Running	0	Feb 19, 2018, 12:58:08 PM
1	qkimage1-5694c94cfd-girpc	Running	0	Feb 19, 2018, 1:18:20 PM
1	qkimage1-5694c94cfd-q59kh	Running	0	Feb 19, 2018, 1:18:20 PM

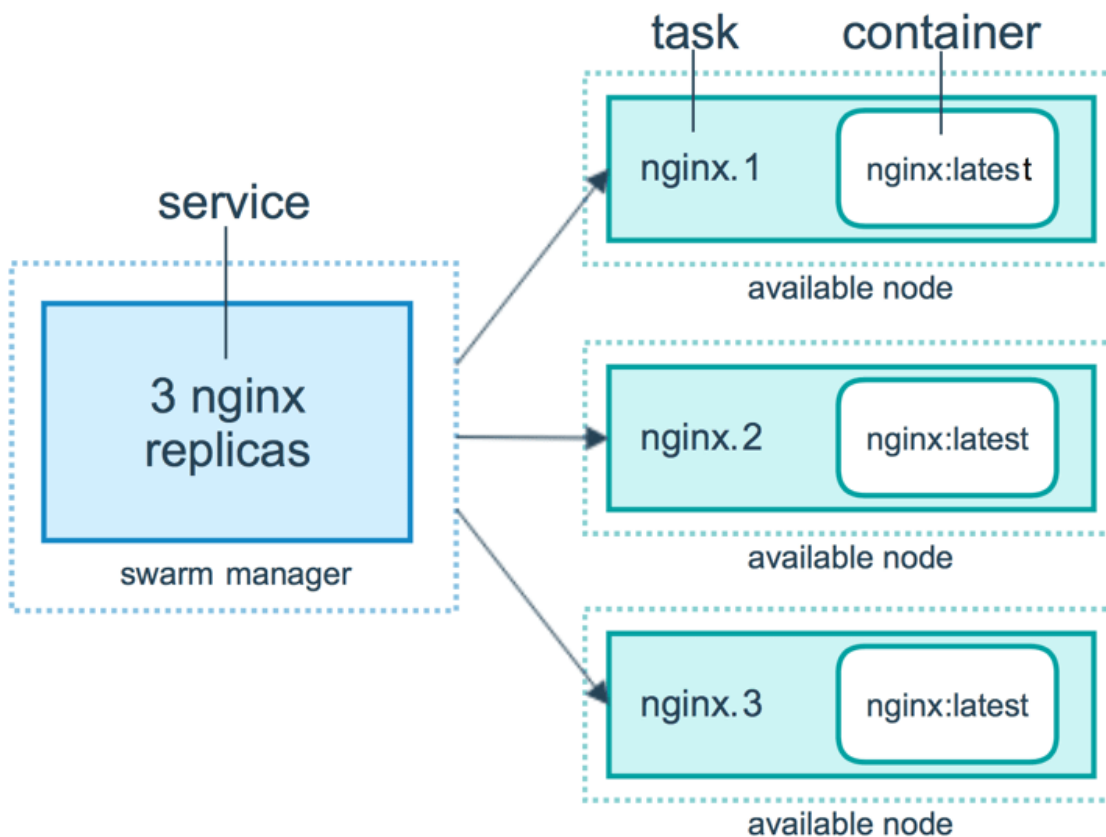
Now the service will be available in all the three pods which will be running the service qkimage1 and image quickstart-image.

## Docker Swarm

Docker Swarm is a clustering and scheduling tool for Docker containers. With Swarm, IT administrators and developers can establish and manage a cluster of Docker nodes as a single virtual system.

Swarms are a cluster of nodes which comprise of:

- **Manager nodes:** Control orchestration, cluster management, and the distribution of tasks.
- **Worker nodes:** The sole purpose of workers is to run containers and services as assigned by a manager node.
- **Services:** A service describes how you'd like an individual container to distribute itself across your nodes. To create a service, specify the exact information as in an ordinary 'docker run', plus new parameters (i.e., # of container replicas).
- **Tasks:** Single containers place work within these "slots" according to the Swarm manager.



## Steps to create a swarm :

To create a set of Docker machines that will act as nodes in our Docker Swarm.

For creating Manager Node:

*docker-machine create --driver hyperv manager1* //For OS with Hyper-V manager

Or

*docker-machine create --driver virtualbox manager1* //For Docker Toolbox

```
Akash@DESKTOP-BBTTUHF MINGW64 ~
$ docker-machine create --driver virtualbox manager1
Running pre-create checks...
Creating machine...
(manager1) Copying C:\Users\Akash\.docker\machine\cache\boot2docker.iso to C:\Users\Akash\.docker\machine\machines\manager1\boot2docker.iso...
(manager1) Creating VirtualBox VM...
(manager1) Creating SSH key...
(manager1) Starting the VM...
(manager1) Check network to re-create if needed...
(manager1) Windows might ask for the permission to configure a dhcp server. Sometimes, such confirmation window is minimized in the taskbar.
(manager1) Waiting for an IP...
Waiting for machine to be running, this may take a few minutes...
Detecting operating system of created instance...
Waiting for SSH to be available...
Detecting the provisioner...
Provisioning with boot2docker...
Copying certs to the local machine directory...
Copying certs to the remote machine...
Setting Docker configuration on the remote daemon...
Checking connection to Docker...
Docker is up and running!
To see how to connect your Docker Client to the Docker Engine running on this virtual machine, run: C:\Program Files\Docker Toolbox\docker-machine.exe env manager1

Akash@DESKTOP-BBTTUHF MINGW64 ~
```

For creating Worker Nodes:

*docker-machine create --driver virtualbox worker1*

```
Akash@DESKTOP-BBTTUHF MINGW64 ~
$ docker-machine create --driver virtualbox worker1
Running pre-create checks...
Creating machine...
(worker1) Copying C:\Users\Akash\.docker\machine\cache\boot2docker.iso to C:\Users\Akash\.docker\machine\machines\worker1\boot2docker.iso...
(worker1) Creating VirtualBox VM...
(worker1) Creating SSH key...
(worker1) Starting the VM...
(worker1) Check network to re-create if needed...
(worker1) Windows might ask for the permission to configure a dhcp server. Sometimes, such confirmation window is minimized in the taskbar.
(worker1) Waiting for an IP...
Waiting for machine to be running, this may take a few minutes...
Detecting operating system of created instance...
Waiting for SSH to be available...
Detecting the provisioner...
Provisioning with boot2docker...
Copying certs to the local machine directory...
Copying certs to the remote machine...
Setting Docker configuration on the remote daemon...
Checking connection to Docker...
Docker is up and running!
To see how to connect your Docker Client to the Docker Engine running on this virtual machine, run: C:\Program Files\Docker Toolbox\docker-machine.exe env worker1
```

```
docker-machine create --driver virtualbox worker2
```

```

C:\Users\Akash>docker-machine create --driver virtualbox worker2
$ docker-machine create --driver virtualbox worker2
Running pre-create checks...
Creating machine...
(worker2) Copying C:\Users\Akash\.docker\machine\cache\boot2docker.iso to C:\Users\Akash\.docker\machine\machines\worker2\boot2docker.iso...
(worker2) Creating VirtualBox VM...
(worker2) Creating SSH key...
(worker2) Starting the VM...
(worker2) Check network to re-create if needed...
(worker2) Windows might ask for the permission to configure a dhcp server. Sometimes, such confirmation window is minimized in the taskbar.
(worker2) Waiting for an IP...
Waiting for machine to be running, this may take a few minutes...
Detecting operating system of created instance...
Waiting for SSH to be available...
Detecting the provisioner...
Provisioning with boot2docker...
Copying certs to the local machine directory...
Copying certs to the remote machine...
Setting Docker configuration on the remote daemon...
Checking connection to Docker...
Docker is up and running!
To see how to connect your Docker Client to the Docker Engine running on this virtual machine, run: C:\Program Files\Docker Toolbox\docker-machine.exe env worker2

```

```
docker-machine ls //to get list of docker machines
```

```

$ docker-machine ls

```

NAME	ACTIVE	DRIVER	STATE	URL	SWARM	DOCKER	ERRORS
default	*	virtualbox	Running	tcp://192.168.99.100:2376		v18.02.0-ce	
manager1	-	virtualbox	Running	tcp://192.168.99.101:2376		v18.02.0-ce	
worker1	-	virtualbox	Running	tcp://192.168.99.102:2376		v18.02.0-ce	
worker2	-	virtualbox	Running	tcp://192.168.99.103:2376		v18.02.0-ce	

## Ip's of docker machines.

```
Akash@DESKTOP-BBTTUHF MINGW64 ~  
$ docker-machine ip manager1  
192.168.99.101  
  
Akash@DESKTOP-BBTTUHF MINGW64 ~  
$ docker-machine ip worker1  
192.168.99.102  
  
Akash@DESKTOP-BBTTUHF MINGW64 ~  
$ docker-machine ip worker2  
192.168.99.103
```

**To SSH into any of the machines.**

```
docker-machine ssh <machine-name>
```

*docker-machine ssh manager1*

[illegible]

## To set up the Swarm.

*docker swarm init --advertise-addr MANAGER\_IP*

```
docker@manager1:~$ docker swarm init --advertise-addr 192.168.99.101
Swarm initialized: current node (vod2qiqcocjesh62dohqopxep) is now a manager.

To add a worker to this swarm, run the following command:

    docker swarm join --token SWMTKN-1-0zz0o252oa1ajaezixm3ggnb8h4noxcmx4yo561f7w2av85dje-99fqm8zp1ykm5tpw74esypxlq 192.168.99.101:2377

To add a manager to this swarm, run 'docker swarm join-token manager' and follow the instructions.
```

**To find out what docker swarm command to use to join as a node, you will need to use the `join-token <role>` command.**

*docker swarm join-token worker*

```
docker@manager1:~$ docker swarm join-token worker
To add a worker to this swarm, run the following command:

    docker swarm join --token SWMTKN-1-0zz0o252oa1ajaezixm3ggnb8h4noxcmx4yo561f7w2av85dje-99fqm8zp1ykm5tpw74esypxlq 192.168.99.101:2377
```

*docker swarm join-token manager*

```
docker@manager1:~$ docker swarm join-token manager
To add a manager to this swarm, run the following command:

    docker swarm join --token SWMTKN-1-0zz0o252oa1ajaezixm3ggnb8h4noxcmx4yo561f7w2av85dje-aysng5gsdsyjl1hpw1y8uzfwcr 192.168.99.101:2377
```

**To get the list of nodes in the swarm.**

*docker node ls*

```
docker@manager1:~$ docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
vod2qiqcocjesh62dohqopxep *	manager1	Ready	Active	Leader

**SSH into the worker1 machine.**

Open a new docker terminal.

*docker-machine ssh worker1*

```
Akash@DESKTOP-BBTTUHF MINGW64 ~
$ docker-machine ssh worker1

##
## ##
## ## ## ## ##
/ ~~~~~ \
~~~~~ /
  o
  \
  /
  \
  /
  \

boot2docker

Boot2Docker version 18.02.0-ce, build HEAD : 99245f4 - Thu Feb  8 17:43:39 UTC 2018
Docker version 18.02.0-ce, build fc4de44
docker@worker1:~$ docker swarm join-token worker
Error response from daemon: This node is not a swarm manager. Use "docker swarm init" or "docker swarm join" to connect this node to swarm and try again.
```

**To join worker1 to the swarm.**

Use the join-token generated for worker

```
docker@worker1:~$ docker swarm join --token SWMTKN-1-0zz0z52oa1ajaezi3xm3ggnb8h4noxcw4y0561f7w2av85d4je-99fqm8zpiylkm5tpw74esyypxlq 192.168.99.101:2377
This node joined a swarm as a worker.
docker@worker1:~$
```

```
docker node ls           //on the manager1
```

```
docker@manager1:~$ docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
vod2qiqcocjesh62dohqopxep *	manager1	Ready	Active	Leader
ufh55mibj423fnikw3fzc3nse	worker1	Ready	Active	

## SSH into the worker2 machine.

Open a new docker terminal.

```
docker-machine ssh worker2
```

[illegible]

## To join worker2 to the swarm

Use the join-token generated for worker

```
docker@worker2:~$ docker swarm join --token SWMTKN-1-0zz0o252oa1ajaeizxm3ggnb84noxcwx4yo561f7w2av85dje-99fqm8zpl1ym5tpw74esypxlg 192.168.99.101:2377
This node joined a swarm as a worker.
docker@worker2:~$
```

```
docker node ls           //on the manager1
```

```
docker@manager1:~$ docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
vod2qiqcocjesh62dohqopxep *	manager1	Ready	Active	Leader
ufh55mibj423fnikw3fzc3nse	worker1	Ready	Active	
9uzkaso7yb9ttlq9cuj0bae6u	worker2	Ready	Active	

```
docker@manager1:~$
```

## To get swarm information.

*docker info* //on the manager1 and zoom into the swarm section

```
docker@manager1:~$ docker info
Containers: 0
  Running: 0
  Paused: 0
  Stopped: 0
Images: 0
Server Version: 18.02.0-ce
Storage Driver: aufs
  Root Dir: /mnt/sda1/var/lib/docker/aufs
  Backing Filesystem: extfs
  Dirs: 0
  Dirperm1 Supported: true
Logging Driver: json-file
Cgroup Driver: cgroupfs
Plugins:
  Volume: local
  Network: bridge host macvlan null overlay
  Log: awslogs fluentd gcplogs gelf journald json-file logentries splunk syslog
Swarm: active
  NodeID: vod2qiqcocjesh62dohqopxep
  Is Manager: true
  ClusterID: r0j7h9o9uieil5qn024vm7yy0
  Managers: 1
  Nodes: 3
  Orchestration:
    Task History Retention Limit: 5
  Raft:
    Snapshot Interval: 10000
    Number of Old Snapshots to Retain: 0
    Heartbeat Tick: 1
    Election Tick: 3
  Dispatcher:
    Heartbeat Period: 5 seconds
  CA Configuration:
    Expiry Duration: 3 months
    Force Rotate: 0
  Autolock Managers: false
  Root Rotation In Progress: false
  Node Address: 192.168.99.101
  Manager Addresses:
    192.168.99.101:2377
```

## To create a Service.

```
docker@manager1:~$ docker service create --replicas 2 -p 8080:80 --name web nginx
uma8aodbr2bbb4wy8uf2zjqis
overall progress: 2 out of 2 tasks
1/2: running [=====>]
2/2: running [=====>]
verify: Service converged
docker@manager1:~$
```

## Check status of service.

```
docker@manager1:~$ docker service ls
ID                NAME      MODE      REPLICAS  IMAGE      PORTS
uma8aodbr2bb     web       replicated 2/2        nginx:latest *:8080->80/tcp

docker@manager1:~$ docker service ps web
ID                NAME      IMAGE      NODE      DESIRED STATE  CURRENT STATE      ERROR      PORTS
9tepjllqjzft     web.1     nginx:latest  worker1   Running        Running 3 minutes ago
r10pvefo1b9t     web.2     nginx:latest  manager1   Running        Running 3 minutes ago

docker@manager1:~$ docker ps
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS      NAMES
a63f21833ba7  nginx:latest  "nginx -g 'daemon of..."  5 minutes ago  Up 5 minutes  80/tcp     web.2.r10pvefo1b9tygq86jotqepre

docker@manager1:~$
```

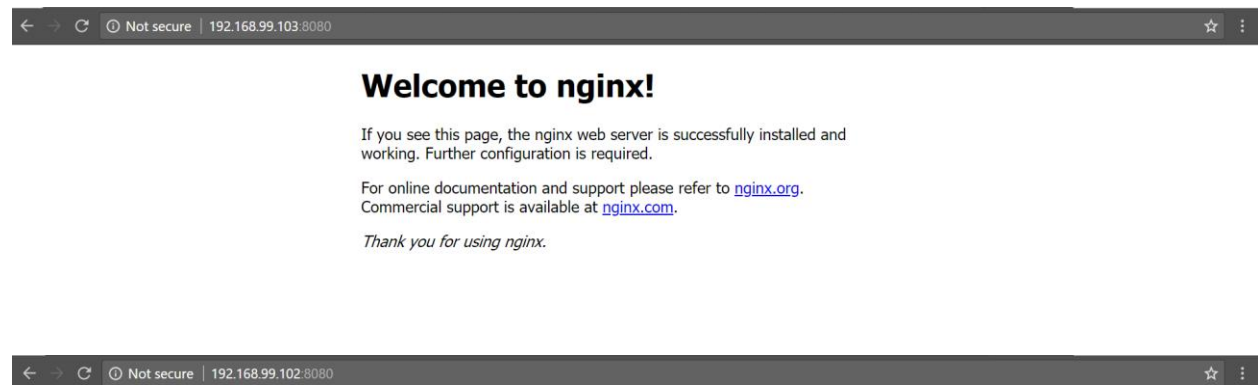
list of docker machines and their ip's //on new docker terminal

```
Akash@DESKTOP-BBTTUHF MINGW64 ~
$ docker-machine ls
NAME      ACTIVE   DRIVER      STATE     URL                  SWARM   DOCKER      ERRORS
default  *        virtualbox  Running  tcp://192.168.99.100:2376   v18.02.0-ce
manager1  -        virtualbox  Running  tcp://192.168.99.101:2376   v18.02.0-ce
worker1   -        virtualbox  Running  tcp://192.168.99.102:2376   v18.02.0-ce
worker2   -        virtualbox  Running  tcp://192.168.99.103:2376   v18.02.0-ce
```

## To access the Service.

http://<machine ip>:8080 in the browser.

You should be able to get the standard NGINX Home page.



## Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to [nginx.org](http://nginx.org).  
Commercial support is available at [nginx.com](http://nginx.com).

*Thank you for using nginx.*





## Scaling up.

To Scale up to 5 replicas from 2 replicas

```
docker@manager1:~$ docker service scale web=5
web scaled to 5
overall progress: 5 out of 5 tasks
1/5: running [=====>]
2/5: running [=====>]
3/5: running [=====>]
4/5: running [=====>]
5/5: running [=====>]
verify: Service converged
docker@manager1:~$
```

To check status of service

```
docker@manager1:~$ docker service ls
ID            NAME      MODE          REPLICAS  IMAGE      PORTS
uma8aodbr2bb  web       replicated    5/5        nginx:latest *:8080->80/tcp

docker@manager1:~$ docker service ps web
ID            NAME      IMAGE          NODE        DESIRED STATE  CURRENT STATE      ERROR          PORTS
9tepjllqjzft  web.1     nginx:latest   worker1     Running        Running 20 minutes ago
r10pvefo1b9t  web.2     nginx:latest   manager1    Running        Running 20 minutes ago
njcihgnd39i   web.3     nginx:latest   worker1     Running        Running about a minute ago
oo4modmbzbb   web.4     nginx:latest   worker2     Running        Running about a minute ago
7ag1uaqj8p48  web.5     nginx:latest   worker2     Running        Running about a minute ago

docker@manager1:~$ docker ps
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS          NAMES
a63f21833ba7  nginx:latest  "nginx -g 'daemon of..."  20 minutes ago  Up 20 minutes  80/tcp         web.2.r10pvefo1b9tygq86jotqepre

docker@manager1:~$
```

Checking tasks running on manager1

```
docker@manager1:~$ docker node ps self
ID            NAME      IMAGE          NODE        DESIRED STATE  CURRENT STATE      ERROR          PORTS
r10pvefo1b9t  web.2     nginx:latest   manager1    Running        Running 23 minutes ago

docker@manager1:~$
```

Checking tasks running worker1 and worker2

```
docker@manager1:~$ docker node ps worker1
ID            NAME      IMAGE          NODE        DESIRED STATE  CURRENT STATE      ERROR          PORTS
9tepjllqjzft  web.1     nginx:latest   worker1     Running        Running 24 minutes ago
njcihgnd39i   web.3     nginx:latest   worker1     Running        Running 5 minutes ago

docker@manager1:~$ docker node ps worker2
ID            NAME      IMAGE          NODE        DESIRED STATE  CURRENT STATE      ERROR          PORTS
oo4modmbzbb   web.4     nginx:latest   worker2     Running        Running 5 minutes ago
7ag1uaqj8p48  web.5     nginx:latest   worker2     Running        Running 5 minutes ago

docker@manager1:~$
```

## Inspecting a node.

docker node inspect [--pretty] Nodename

```
docker@manager1:~$ docker node inspect --pretty worker1
ID: ufh55mbj423fnikw3fzc3nse
Hostname: worker1
Joined at: 2018-02-14 06:05:07.318961913 +0000 utc
Status:
  State: Ready
  Availability: Active
  Address: 192.168.99.102
Platform:
  Operating System: linux
  Architecture: x86_64
Resources:
  CPUs: 1
  Memory: 995.9MiB
Plugins:
  Log: awslogs, fluentd, gcplogs, gelf, journald, json-file, logentries, splunk, syslog
  Network: bridge, host, macvlan, null, overlay
  Volume: local
Engine Version: 18.02.0-ce
Engine Labels:
  - provider=virtualbox
TLS Info:
  TrustRoot:
  -----BEGIN CERTIFICATE-----
  MITIBajCCARCGAwIBAgIUda+SP/yFFCxEicgVQnEP1oKCfSowCgYIKoZIzj0EAwIw
  EZERMA8GA1UEAxMIc3dhcm0tY2EwHhcNMjE0MDU0MzAwHhcNMjE0MDU0
  MzAwHjATMREuDuYDVQDEwHdz2Fybs1jYTBZMBMGByqGSM49AgEGCCqGSM49AwEH
  A0IABJ5kqJhb8xsa0D+dzX1DxjYm8q5xHR6J2topEiPLHF9qZwmFFBuigI9nFICz
  MNemGoTxy/U+vo/3SjsCwM1xSAiJQjBAMA4GA1UdDwEB/wQEAwIBBjAPBgNVHRMB
  Af8EBTADAQH/MB0GA1UdDgQqMBQKkr9P4tgKcggNBk517LJJ+e+0QDAKBggqhkJ0
  PQQDAgNIADBFaiEA4FQaxsUdjiUCQs8170eD1ZJExtRSJUDtSI4hw7sZZnQCIFhT
  YdEewMrWHDic7OLRBBMSne5JTjFqbucRnrDbHs11
  -----END CERTIFICATE-----
  Issuer Subject: MBMxETAPBgNVBAMTCHN3YXJtLWln
  Issuer Public Key: MFkwEwYHKOZIzj0CAQYIKoZIzj0DAQcDQgAEnmSomFvzGxrnQP53NFUPGNgzrnmEdHona2ikQiksd/2pnCZ98G6KAj2d8gLMw14wahPHL9T6+j/dKOwJYyXFICA==
```

## Draining a node.

docker node update --availability drain Nodename

```
docker@manager1:~$ docker node update --availability drain worker1
worker1
docker@manager1:~$ docker service ps web
ID NAME IMAGE NODE DESIRED STATE CURRENT STATE ERROR PORTS
orasm2f1bwjz web.1 nginx:latest manager1 Running Running 17 seconds ago ERROR PORTS
9tepjl1qjzft \_ web.1 nginx:latest worker1 Shutdown Shutdown 16 seconds ago
r10pvefo1b9t web.2 nginx:latest manager1 Running Running 35 minutes ago
x6x76g9h7kos web.3 nginx:latest manager1 Running Running 17 seconds ago
njcihgnd39i \_ web.3 nginx:latest worker1 Shutdown Shutdown 16 seconds ago
oo4modmbezbb web.4 nginx:latest worker2 Running Running 16 minutes ago
7ag1uaqj8p48 web.5 nginx:latest worker2 Running Running 16 minutes ago
```

## Checking replicas running on worker1,worker2 and manager1

```
docker@manager1:~$ docker node ps worker1
ID NAME IMAGE NODE DESIRED STATE CURRENT STATE ERROR PORTS
9tepjl1qjzft web.1 nginx:latest worker1 Shutdown Shutdown 47 seconds ago
njcihgnd39i web.3 nginx:latest worker1 Shutdown Shutdown 47 seconds ago
docker@manager1:~$ docker node ps worker2
ID NAME IMAGE NODE DESIRED STATE CURRENT STATE ERROR PORTS
oo4modmbezbb web.4 nginx:latest worker2 Running Running 16 minutes ago
7ag1uaqj8p48 web.5 nginx:latest worker2 Running Running 16 minutes ago
docker@manager1:~$ docker node ps manager1
ID NAME IMAGE NODE DESIRED STATE CURRENT STATE ERROR PORTS
orasm2f1bwjz web.1 nginx:latest manager1 Running Running about a minute ago
r10pvefo1b9t web.2 nginx:latest manager1 Running Running 35 minutes ago
x6x76g9h7kos web.3 nginx:latest manager1 Running Running about a minute ago
docker@manager1:~$
```

## Inspecting the Service.

```
docker@manager1:~$ docker service inspect --pretty web

ID:                uma8aodbr2bbb4wy8uf2zjq18
Name:              web
Service Mode:     Replicated
  Replicas:        5
Placement:
UpdateConfig:
  Parallelism:     1
  On failure:      pause
  Monitoring Period: 5s
  Max failure ratio: 0
  Update order:    stop-first
RollbackConfig:
  Parallelism:     1
  On failure:      pause
  Monitoring Period: 5s
  Max failure ratio: 0
  Rollback order:  stop-first
ContainerSpec:
  Image:           nginx:latest@sha256:285b49d42c703fdf257d1e2422765c4ba9d3e37768d6ea83d7fe2043dad6e63d
Resources:
Endpoint Mode:    vip
Ports:
  PublishedPort = 8080
  Protocol      = tcp
  TargetPort    = 80
  PublishMode   = ingress
```

## Removing the Service.

`docker service rm ServiceName`

```
docker@manager1:~$ docker service rm web
web
docker@manager1:~$ docker service ls
ID                NAME                MODE                REPLICAS            IMAGE                PORTS
docker@manager1:~$ docker service inspect --pretty web
Error: no such service: web
docker@manager1:~$ docker node ps worker1
ID                NAME                IMAGE                NODE                DESIRED STATE        CURRENT STATE        ERROR                PORTS
docker@manager1:~$ docker node ps worker2
ID                NAME                IMAGE                NODE                DESIRED STATE        CURRENT STATE        ERROR                PORTS
docker@manager1:~$ docker node ps manager1
ID                NAME                IMAGE                NODE                DESIRED STATE        CURRENT STATE        ERROR                PORTS
docker@manager1:~$ docker service ps web
no such service: web
docker@manager1:~$
```

## Stopping docker machine.

```
Akash@DESKTOP-BBTTUHF MINGW64 ~  
$ docker-machine ls  
NAME      ACTIVE  DRIVER      STATE      URL                      SWARM   DOCKER      ERRORS  
default   *       virtualbox   Running    tcp://192.168.99.100:2376   v18.02.0-ce  
manager1  -       virtualbox   Running    tcp://192.168.99.101:2376   v18.02.0-ce  
worker1   -       virtualbox   Running    tcp://192.168.99.102:2376   v18.02.0-ce  
worker2   -       virtualbox   Running    tcp://192.168.99.103:2376   v18.02.0-ce  
  
Akash@DESKTOP-BBTTUHF MINGW64 ~  
$ docker-machine stop worker1  
Stopping "worker1"...  
Machine "worker1" was stopped.  
  
Akash@DESKTOP-BBTTUHF MINGW64 ~  
$ docker-machine stop worker2  
Stopping "worker2"...  
Machine "worker2" was stopped.  
  
Akash@DESKTOP-BBTTUHF MINGW64 ~  
$ docker-machine stop manager1  
Stopping "manager1"...  
Machine "manager1" was stopped.  
  
Akash@DESKTOP-BBTTUHF MINGW64 ~  
$ docker-machine ls  
bash: docker-machine: command not found  
  
Akash@DESKTOP-BBTTUHF MINGW64 ~  
$ docker-machine ls  
NAME      ACTIVE  DRIVER      STATE      URL                      SWARM   DOCKER      ERRORS  
default   *       virtualbox   Running    tcp://192.168.99.100:2376   v18.02.0-ce  
manager1  -       virtualbox   Stopped    tcp://192.168.99.101:2376   Unknown  
worker1   -       virtualbox   Stopped    tcp://192.168.99.102:2376   Unknown  
worker2   -       virtualbox   Stopped    tcp://192.168.99.103:2376   Unknown
```

## Deleting docker-machine.

```
Akash@DESKTOP-BBTTUHF MINGW64 ~  
$ docker-machine rm worker1 worker2 manager1  
About to remove worker1, worker2, manager1  
WARNING: This action will delete both local reference and remote instance.  
Are you sure? (y/n): y  
Successfully removed worker1  
Successfully removed worker2  
Successfully removed manager1  
  
Akash@DESKTOP-BBTTUHF MINGW64 ~  
$ docker-machine ls  
NAME      ACTIVE  DRIVER      STATE      URL                      SWARM   DOCKER      ERRORS  
default   *       virtualbox   Running    tcp://192.168.99.100:2376   v18.02.0-ce  
  
Akash@DESKTOP-BBTTUHF MINGW64 ~  
$
```

## **Kubernetes Vs Docker Swarm.**

<https://platform9.com/blog/compare-kubernetes-vs-docker-swarm/>

## Microservices

- Software systems become more complex as their scale
  - scope
  - Volume
  - user interactions
- Sometimes system that grow in size beyond the boundaries initially define pose a particular problems when it comes to changing them

### What is Microservices?

- Microservice architecture are used to achieve faster delivery and greater safety as the scale of their systems increase
- Microservices are goal-oriented
- Microservices are focused on replaceability
- This idea that driving toward replacement of components rather than maintaining existing components get to the very heart of what makes the microservices approach special

## Microservices

- More specifically, the real value of microservices is realized when we focus on two key aspects - **speed and safety**
- **The Speed of Change** - The desire for speed is a desire for immediate change and ultimately a desire for adaptability
- **At Scale** - Systems that can work at scale don't break when under pressure; instead they incorporate built-in mechanisms to increase capacity in a safe way
- **In Harmony** - Organizations that succeed with microservice architecture are able to maintain their system stability while increasing their change velocity. In other words, they created a harmony of speed and safety that works for their own context

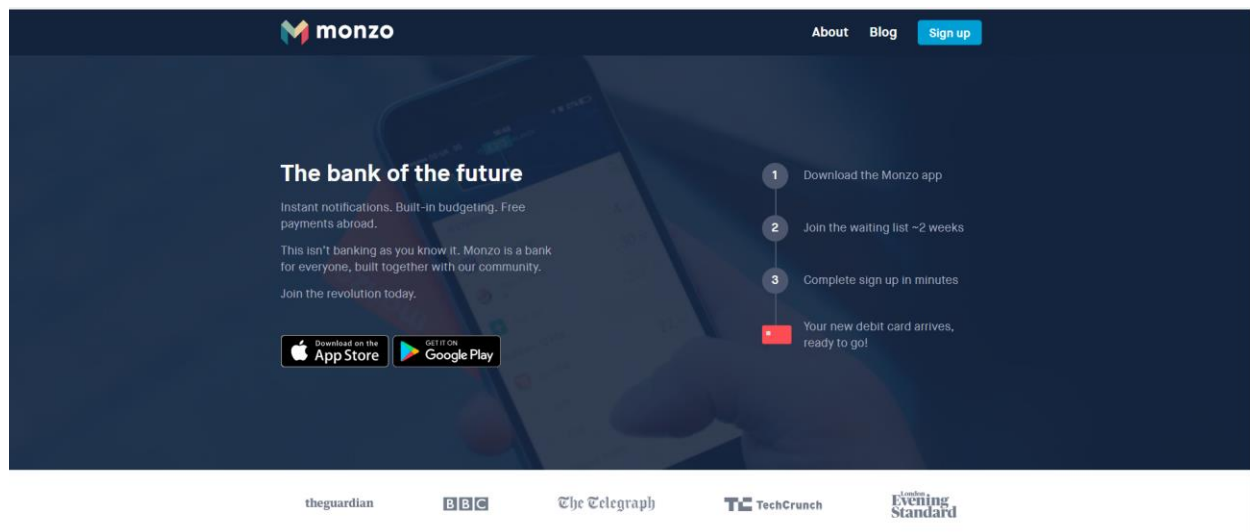
## How it impact Business ?

- **Agility** allows organizations to deliver new products, functions, and features more quickly and pivot more easily if needed
- **Comprehensibility of the software system** simplifies development planning, increases accuracy, and allows new resources to come up to speed more quickly
- **Independent deployment** of components gets new features into production more quickly and provides more flexible options for piloting and prototyping
- Organizational alignment of services to teams reduces ramp-up time and encourages teams to **build more complex products** and features iteratively
- **Polyglotism** permits the use of the right tools for the right task, thus accelerating technology introduction and increasing solution options

## Case Studies

### MONZO: Build a Modern Bank Backend.

Ref: <https://monzo.com/blog/2016/09/19/building-a-modern-bank-backend/>



**Objective:** At Monzo, we're building a banking system from scratch. Our systems must be available 24x7 no matter what happens, scalable to hundreds of millions of customers around the world, and very extensible. This first post in a series about our platform explains how we're building systems to meet these demands using modern, open-source technology.

**Strategy:** Building the backend as a collection of distributed **microservices**. For an early-stage startup, this architecture is quite unusual; most companies start with a centralised application using well-known frameworks and relational databases.



But with an ambition to build the best current account in the world and scale it to hundreds of millions of customers, we knew that we had to be the bank with the best technology.

**Implementation:** When we launched the Monzo **beta**, our backend had grown to nearly 100 services (now about 150), and our engineering team had grown considerably too. Few areas of focus: Cluster management, polyglot services, RPC transport, Asynchronous messages.

After a year or so of using Mesos and Marathon, we decided to switch to **Kubernetes**, which we run on a fleet of **CoreOS** machines. It is very tightly integrated with Docker, and is the product of Google's **long experience** running containers in production at global scale. Deploying Kubernetes in a **highly available** configuration on AWS is not for the faint of heart and requires you to get familiar with its internals, but we are very pleased with the results. We regularly kill hosts in production and Kubernetes quickly reschedules applications to accommodate.

There are increasingly large number of such startups and large corporation wanting to implement microservice strategy in developing their product through technology. This advancement can improve their efficiency and transactions can be monitored easily.

**Conclusion:** Today Fintech is rising at an alarming pace with huge investment from major companies competing for the best technology for driving their business needs and goals. Kubernetes and Docker Swarm are the tools and techniques on how systems operate in a micro service manner.